Understanding & Analyzing the Garbage-First Family of Garbage Collectors

Wenyu Zhao

© Wenyu Zhao 2018

Except where otherwise indicated, this thesis is my own original work.

Wenyu Zhao
21 October 2018

To my partner and my parents.

Acknowledgments

Who do you want to thank?

Abstract

Region-based garbage collectors, including Garbage-first GC, Shenandoah GC, C4 GC and ZGC, share a lot of common designs and algorithms, such as the similar marking algorithm, memory structure and allocation algorithm. However, the existence of such underlying relationship among these collectors has never been well identified and explored. Instead, the designers of these region-based collectors tend to treat these as stand-alone collectors instead of an improvement to the existing collector. Such ignorance of the underlying relationships among the region-based collectors can mislead the future design of related garbage collectors.

Hence, analysis, measurements and comparisons among these collectors can also be hard. Although the design and algorithms of these collectors are similar, their structural relationships are not reflected to the original design and implementations of these collectors. For this reason no one have ever succsessfully measured the GC performance contribution of some specific part of the GC algorithm or the extension component involved in these collectors. Which means no one can properly understand the pros and cons of the design of these collectors, and may further leads to some potential performance issue due to the inappropriate GC design.

As the fundamental novel contribution, this thesis is the first to prove the existence of a potential structural relationship among the region-based collectors and makes a deep exploration of the structural relationships among a set of region-based collectors which have similar design to the Garbage-First GC. In thesis I use the term "The Garbage-First Family of Garbage Collectors" to describe such category of collectors.

This thesis produces the first implementation to reflect the previously discovered structural relationships. Sperifically, in this thesis, I discuss the implementation details of a total of six G1 family of collectors, starting from a simple region-based collector to the Garbage-first GC and Shenandoah GC. Each collector are implemented as a improved version of the previous collector to reflected the corresponding algorithmic relationship.

Based on such implementation, this thesis performs a detailed and careful analysis of the GC performance contribution of each component of the algorithms. This includes the measurement of the GC pause time, several mutator barrier overheads and the space overhead of remembered-sets.

The exploration of the Garbage-First Family of Garbage Collectors leads to a conclusion that there exist a structural relationship among these G1 family of collectors. Instead of being stand-alone collector, they tend to be collectors with algorithmic improvements to some existing collectors.

As the result of the GC performance evaluation, structural components including concurrent-marking, remembered-sets and concurrent evacuation contributes to a improvement of ???%, ???% and ???% respectively to the average GC pause time

on the DaCapo benchmark suite. By using remembered-sets, G1 has ???% average footprint overhead. Mutator barriers performance SATB barrier, remembered-set barrier, and Brooks barrier increases the mutator overheads by ???%, ???% and ???% respectively.

The explored algorithmic relationships among the G1 family of collectors can help GC designers or other programmers working with JVM to have a deeper understanding of the G1 family of collectors as well as their relationships and the pros and cons of their involved improvements. The measured GC performance contribution of each component of the GC algorithm can help garbage collection algorithm designers to reconsider the design of the region-based garbage collectors and memory structures, identify the main advantages and drawbacks of each involved GC algorithm and hence have the ability to make further optimizations to them.

Contents

List of Figures

List of Tables

Introduction

This thesis explores the Garbage-First family of garbage collectors, including the underlying relationship among them, the improvements each GC does and the performance contribution by each algorithmic component. My thesis is that different Garbage-First family of collectors are strongly related and each collector is actually an improvements to other existing collectors. These improvements increases the GC performance in some aspect but also cause several extra drawbacks.

1.1 Project Statement

As the memory size of modern server machines becomes larger, latency time of garbage collectors for managed programming languages generally becomes longer. In this way, the designing of low-latency garbage collectors has become a hot topic today. This involves reducing the latency of collectors and performing heap compaction or evacuation to avoid heap fragmentation.

Among all the existing low-latency garbage collectors, the category of region-based collectors is widely explored and used in modern industry. Region-based collectors are designed to reach high GC performance (especially GC pause times) by managing the heap as a set of memory regions. As four well-known region-based garbage collectors, Garbage-first GC, Shenandoah GC, C4 GC and ZGC created by [Detlefs et al., 2004; Flood et al., 2016; Tene et al., 2011; Liden and Karlsson, 2018] respectively are high performance region-based garbage collectors implemented for Java virtual machines to reach short GC pause times and high program throughput.

The design and implementation of such region-based collectors has been well explored nowadays, but unfortunately their underlying relationship is still remains unidentified and unexplored. Although these collectors share a lot of basic algorithms and structures, the original papers or other publications of these collectors only regard them as newly created and stand-alone collectors instead of improvements to the existing collectors. Such ignorance of the potential relationships prevents GC designers to have a clear overview of the algorithms supporting these collectors.

Also, the lack of explorations on these relationships can make the GC performance evaluation and analysis among these collectors to be impossible. The existing G1 family of collectors, including G1 GC, Shenandoah GC, C4 GC and ZGC are invented

by different groups or campanies (Sun, RadHat, Azul and Oracle respectively). The implementations and optimizations on these collectors are vary. This makes the original implementations of the G1 family of collectors do not reflect their underlying structural relationships. The implementation differences among these collectors make them become impossible to be carefully measured for the GC performance impact of each structural component. This means no one have ever truely understand the GC performance contribution of each GC algorithm involved in these collectors. On the one hand such lack of explorations on these relationships and the lack of performance analysis and comparisons may cause some unexpected performance issue to these collectors. As an example, according to BrianDemsky [2015], sometimes GC time of region-based collectors can take up to 50% of the total execution time of Big Data systems. On the other hand this can also mislead the future design of related collectors.

So here come to the problem which this thesis is aiming to solve: What are the relationships among the G1 family of collectors and in which way will they contribute to the GC performance?

1.2 Contribution

I insist that the G1 family of collectors are strongly related to each other, although the original papers of these collectors present them as a stand-alone collector instead of a improved collector based on other existing GC algorithms. The lack of explorations to such structural relationships makes it hard to carefully analyze the performance contribution caused by each part of the GC algorithm, hence prevents GC designers to truly understand the pros and cons of each GC algorithms.

This thesis makes a deep exploration to the underlying relationship among the G1 family of garbage collectors. The algorithmic relationships are summarised as a series of progressive improvements, which generally reflects the evolution history of the G1 family of collectors.

As another major component of the exploration, I implement the collectors by following the steps of the progressive improvements, staring from a most simple form of the region-based collector to the most complex collector, e.g. Garbage-First GC and Shenandoah GC, to reveal the hierarchy of the G1 family of collectors. GC performance of these implemented collectors are also evaluated and discussed.

The general steps involved in thie thesis for exploring these collectors includes:

Discover the relationships of the G1 family of garbage collectors includs understand the basic algorithms of these collectors, infer the underlying relationships among them and reconsider them as a series of algorithm improvements instead of individual collectors.

Implement a simple mark-region collector by starting from an existing SemiSpace GC in MMTk and replace the two copy-spaces with a region-space which divides the memory up into multiple regions. Then adding an extra marking phase before the evacuation phase.

Perform a series of improvements by starting from the simple mark-region collector and progressively perform improvements including linear-scan evacuation, concurrent marking, remembered sets and concurrent evacuation. This step produces a few improved version of the mark-region collectors, the Garbage-First GC and the Shenandoah GC.

Measure the GC performance includes measuring the GC pause times of all implemented collectors, the barrier overheads and the remembered set size. This measurements reflect the performance gain caused by each algorithmic improvement.

Explain and explore the results of the previously generated measurement results, which leads to a reconsideration of the pros and cons of each G1 related GC algorithms in depth.

A total of six different Garbage-First family of garbage collectors were implemented to reveal the relationships among them and demonstrate different kinds of garbage collection algorithms, such as concurrent marking, remembered-sets and Brooks barriers.

In order to make detailed and careful analysis on these collectors and be able to compare these collectors at an algorithmic level, the implementations are as close as possible to the design of their original papers.

I measure the GC pause times for each implemented collectors, as well as the mutator ooverheads of the mutator barriers involved in these collectors, including snapshot-at-the-beginning (SATB) barrier, remembered-set barrier and the Brooks barrier. I found that concurrent marking contributes to a decrease of the average GC pause time by ???% where linear scanning evacuation increases the average GC pause time by ???%. However, the use of remembered-sets can reduces the GC pause time up to ???%. As the result of barrier latency analysis, I found that SATB barriers can increase the mutator overhead by ???% where Brooks barriers increase the mutator overhead of Shenandoah GC by ???%.

All work were done in JikRVM, a research purpose Java virtual machine and MMTk, a memory management took written in java (Blackburn et al. [2004]).

1.3   Thesis Outline

Chapter 2 provides an general background and an overview of the Garbage-First family of garbage collectors. Similarities and differences among some major categories of garbage collectors are discussed as well. Also the related work on implementing and measuring region-based collector and G1 family of collectors are discussed.

Chapter 3 provides the detailed steps and algorithms of the implementation of all the Garbage-First family of garbage collectors.

Chapter 4 describes the methodology used for evaluating the implemented collectors,

including the benchmark involved and the detailed steps of each evaluation. This chapter also presents the results of the evaluation and benchmarking on the implemented collectors, as well as detailed and critical discussion on the evaluation results of the implemented collectors.

Chapter 5 discusses the work related to this project that I plan to do in the future and the conclusion of this thesis.

Background and Related Work

Garbage collection is a hot topic in terms of modern managed programming language implementations. Specifically, region-based collectors like Garbage-First GC and C4 GC are widely used in modern Java programming languages to archive high GC performance. This chapter describes the background and basic ideas of several garbage collectors, particularly those targeting the Java virtual machine and are implemented in OpenJDK or JikesRVM, as well as the differences among several garbage collection techniques. In addition, this chapter also performs a general discussion of the related work on analyzing region-based garbage collectors.

Section 2.2 roughly discusses and compares different class of GC algorithms. Section 2.3 describes the general design of the Garbage-first family of garbage collectors. Section 2.4 describes the related work on analyzing region-based garbage collectors.

2.1 JikesRVM and MMTk

The whole project discussed in this thesis are based on JikesRVM and all the garbage collectors I implemented and evaluated in this project are all implemented by using MMTk. In this section I will briefly discuss the design of JikesRVM and MMTk, as well as some description on the general structure of them.

2.1.1 JikesRVM

JikesRVM is a research purpose Java virtual machine and was firstly released by Alpern et al. [2005]. It is a meta-circular JVM which is implemented in the Java programming language and is self hosted. JikesRVM was designed to provide a flexible open sourced test-bed to experiment with virtual machine related algorithms and technologies.

In stead of executing Java programs by directly interpreting the Java byte code, JikesRVM compiles them into machine code for execution. JikesRVM implemented two tiers of compilers, the baseline compiler and the optimizing compiler. The baseline compiler simply translates the Java byte code into machine code and do no optimizations while the optimizing compiler performs several optimizations during

the code generation phase. I performed all benchmarking and analysis works on the optimized build of the JikesRVM.

2.1.2   MMTk

MMTk is a memory management toolkit. MMTk is used as a memory management module for JikesRVM and is responsible for memory allocation and garbage collection.

For memory allocation, MMTk defines several address spaces to allocate different type of objects. e.g. `NonMovingSpace` for non-copyable objects and `SmallCodeSpace` for storing java code. After receiving a allocation request, MMTk will decide which space the object belongs to and allocate memory from that space.

All of the G1 family of garbage collectors involved in this thesis involves two major phases, the marking phase and the evecuation phase. Instead of using the MMTk's pre-defined `PREPARE -> SCAN ROOTS -> CLOSURE -> RELEASE` collection phase which only performs a single tracing on the heap, I extended this to perform a seperate full or partial heap tracing or linear scanning phase for evacuation and reference updating.

MMTk will check for stop-the-world or concurrent garbage collection after each requested space allocation. This involves the invocation of methods `collectionRequired(...)` and `concurrentCollectionRequired(...)`. I made full use of these two methods, not only for checking whether a collection is required, but also performs switching between different schedule of collection phases for either nursery or mature collection for the G1 collector.

2.2   Categories of GC Algorithms

This section discusses the main classes of garbage collection algorithms, as well as their pros and cons.

2.2.1   Reference counting

Reference counting is a widely used garbage collection technique which tracks the count of references for each heap allocated object (Detlefs et al. [2002]). The reference count for a object is increased when a new variable references to the object and decreased when a reference to the object is deleted or goes out of its declaring scope. The reference count for a object is initialized to one when the object is created, which means there is only one owner for the object (its creater). When the reference count of the object goes to zero, it is certain that the object has no owner that references to it. Then the object becomes floating garbage and its occupying memory is released.

In order to track the reference count for each object, a write barrier is involved for reference counting collectors. For each object reference modification `obj.x = y`, the reference count of the old object reference is decreased and the reference count of the

new object reference is increased by 1. If the old object reference has no owner, then its memory cell gets swept immediately.

This collection algorithm is highly efficient and the workload for sweeping objects are almost evenly distributed which makes the mutator pause time for reference counting to be extremely short.

However, one major disadvantage of the reference counting GC is that it can hardly handle cyclic references (Lins [1992]) where some object A references the other object B and B also references A. In such case the reference count is at least 1 for both A and B, even if there is no objects referencing to them.

2.2.2 Mark & sweep GC

Mark and Sweep GC is a type of tracing GC which uses the object graph to assist with garbage collection. The algorithm considers all the objects that is the unreachable in the program as garbage (Endo et al. [1997]). In this way Mark and Sweep GC has the ability to collect cyclic referenced garbages, as long as they are unreachable from other live objects.

When allocating objects or requesting memory pages from some memory resource, the mutators dose no extra work but only checks whether the memory is full. If there is no free space for allocation, the execution of all the mutators will be paused and the Mark and Sweep GC is triggered.

The Mark and Sweep GC requires an extra metadata byte in the object header for marking. During each GC cycle, the Mark and Sweep GC firstly scans and marks all the static object, global objects and stack variables as a set of "root objects". Then the collector starting from the root objects and recursively walks over the object graph to mark all the remaining objects. At the end of the marking phase, all the marked objects in the heap are reachable from the "root set" and all other objects become floating garbage and are sweeped.

After a GC cycle, all the paused mutators are resumed to continue execution and allocation.

The Mark and Sweep GC has the ability to collect cyclic referenced garbages by tracing the reachability of heap objects. However, as the use of large-scaled servers and other programs for business becomes more and more porpular, the Mark and Sweep GC reveals its drawback that it can cause significant memory fragmentation after a sufficiently long running time. Because when the collector keeps allocate and free small memory chunks, the size of contiguous free chunks becomes smaller, which may leads to allocation failure for large objects, even if the total free memory size is larger than the requested chunk size.

2.2.3 Copying GC

Copying GC is a class of garbage collectors that aims to reduce heap fragmentation by performing heap compation, which moves objects in the heap together.

As one of the most simple copying GC, the SemiSpace GC divides the whole heap into two spaces, the from-space and the to-space. All objects allocation are done within the from-space. When the from-space is exhausted, a GC cycle is triggered. The collector will starting from all root objects, and recursively walks over the object graph to copy all reachable objects to the to-space. Then the to-space becomes the new from-space for further allocation. In this way the SemiSpace GC ensures all live objects are copied to the to-space and all non-reachable objects (i.e. dead objects) are not forwarded and are swept at the end of the GC cycle.

Copying GCs has the ability to reduce heap fragmentation but can cause longer pause time for each GC cycle. Especially for some GCs which has additional evacuation phase at the end of a marking phase, e.g. the MarkCompact GC.

2.3 Garbage-First Family of Garbage Collectors

This section will give a brief introduction the three most popular Garbage-first family of garbage collectors. For each collector, the basic algorithm and the past and current status will be discussed.

2.3.1 Garbage-First GC

Garbage-First GC is a copying collector which was initially released in Oracle JDK 6 and was fully supported in Oracle JDK 7. G1 was designed as server-style collector which targeting machines with multi-processors and large memories.

G1 GC divides the whole heap up into some fiexed sized regions. As a copying collector, G1 GC tries to reduce the pause time for evacuating objects by performing evacuation on a subset of regions (called the collection set) instead of all allocated regions.

The collection cycle of the Garbage-First GC starts with a concurrent marking phase which the collector threads marks all the objects in the heap, just like the Mark & Sweep GC, but without pausing the mutators. After the marking phase, the G1 collector selects the collection set which contains the regions with smallest ratio of live objects. Then the collector evacuates live objects in the collection set.

In order to perform evacuation on a subset of regions, the collector uses a data structure called "remembered set" to remember all uses of the objects in the collection set. After these live objects are evacuated, the collector scans the remembered set to update the pointers in other regions that references these live objects.

By performing partial heap evacuation, the G1 GC generally has lower pause time than other copying GCs, especially on machines with large heaps (Detlefs et al. [2004]). By adjusting the size of the collection set before evacuation, G1 has the ability to control the pause time to meet some user-defined soft pause time goal. However, one main drawback of G1 is that it is not suitable for small heaps, and the implementation of remembered sets can be inefficient.

2.3.2   Shenandoah GC

Shenandoah GC is concurrently an experimental collector for OpenJDK. Shenandoah GC also divides heap up into regions and performs concurrent marking, similar to the Garbage-first GC.

Shenandoah GC tries to further reduces the GC latency by performing concurrent compaction. The concurrent marking phase that Shenandoah GC has is similar to the G1's concurrent marking phase. However, Shenandoah GC does not have a generational mode and does not perform pertial evacuation to reduce pause times. Instead, the Shenandoah GC performs the evacuation phase concurrently to collect all posible regions, without pausing mutators.

By performing concurrent marking and compaction, Shenandoah GC does most of the heap scanning work in concurrent. In this way the pause time caused by garbage collection is extremely small and is not proportional to the heap size. However, Shenandoah GC has to insert some mutator barriers into the Java program, before every object reference read and write. So the mutator overhead caused by these barrier is much greater than other GCs.

2.3.3   C4 GC

C4 GC is a pauseless GC algorithm created by Tene et al. [2011] and is the default collector of Azul's Zing JVM. C4 is a region-based generational collector which performs mostly concurrent marking and compaction during nursery and mature GCs. A mutator barrier called the Brooks indirection barrier is introduced to support concurrent heap compaction without pausing the mutators. According the Tene et al. [2011], by perform most of the GC work in concurrent, C4 GC has the ability to reduce the mutator response time down to around 10 ms when targeting 100 GB heaps.

2.3.4   ZGC

ZGC is a new garbage collector introduced by Liden and Karlsson [2018] and is very similar to the Shenandoah GC and is also currently under experimental status. ZGC redesigns the mutator barriers of the Shenandoah GC to largely reduces the barrier overhead. ZGC assigns a color for each pointer and stores this metadata into the unused bits in the 64bit pointers. Instead of inserting barriers on every object reference read and write, ZGC only uses a "load barrier" which is only inserted before the mutator loads an object reference from the heap. The barrier is responsible for both object concurrent marking and evacuation, by checking the "color" metadata in the pointer. In this way ZGC remarkably reduces the throughput reduction caused by the mutator barriers.

2.4 Related Work

2.4.1 Implementations and evaluations of the G1 family of collectors

The current working G1 family of garbage collectors implemented for modern Java Virtual Machines includes the Garbage-first GC, Shenandoah GC and ZGC. All of these three major collectors are implemented in OpenJDK.

Detlefs et al. [2004] designed and evaluated the basic algorithm of the Garbage-First collector, including the original design of the pure(non-generational) and generational version of G1 GC. Based on this, G1 GC was firstly implemented and released as an alternative experimental GC for OpenJDK 7. Flood et al. [2016] designed the basic algorithm of the Shenandoah garbage collector. Based on this, Shenandoah GC was firstly implemented as an alternative experimental GC for OpenJDK 8. ZGC is a new garbage collector designed and implemented for OpenJDK and has justly released the the first experimental version in OpenJDK 11.

[Detlefs et al., 2004; Flood et al., 2016] evaluated the performance of G1 GC and Shenandoah GC respectively, including the pause time and mutatur barriers overhead. However, their evacuations uses different benchmarking program and are done within different hardward platform. In addition, the implementations between these collectors varies too much, even for parts that shares the common ideas, which makes it hard to compare the benchmarking results between G1 GC and Shenandoah GC.

2.4.2 Evaluation of region-based collectors

Gay and Aiken [1998] measured the performance of memory management on a region-based heap. They divided the heap up into regions and used a safe C dialect with a conservative reference-counting collector to manage the memory. Then they measured the performance of both allocation and collection over such region-based heap, on a collection of realistic benchmarks. They found that the region structure has advantages for memory allocation and collection, with a competitive performance compared to malloc/free in C programs and has low overhead when using a garbage collector. However, they only evaluated the reference counting collector, which is not a mark and copy collector that this thesis is trying to evacuate. Also, the language they were using is a safe C dialect, which is different to the Java language in many aspects.

Qian and Hendren [2002] implemented a region-based allocator as well as a simple copying collector on JikesRVM. Based on the evaluation results on SPECjvm98 benchmarks, they found that the overhead of allocating and releasing memory pages and the cost of related write barriers is quite low. However, they are not focus too much on evaluating the overhead on a wide range of additional memory management techniques built upon the region-based memory structure, such as concurrent marking or evacuation.

2.4.3   Evaluation of barrier overheads

Yang et al. [2012] evaluated a wide range of different read and write barriers on JikesRVM, using the DaCapo benchmarks and the SPECjvm98 benchmarks. Also they observed the barrier performance differences of in-order and out-of-order machines. They found that write barriers are generally has overheads of 5.4% while write barriers has average overheads of 0.9%. However, these barriers evaluations did not include the evacuation on the Brooks indirection barriers. In addition, this paper measured the behavior of card marking barriers but did not measure the overhead of remembered-set barriers which is used along side with the card marking barrier in the Garbage-first GC.

2.5   Summary

In this section I introduced the background and basic ideas of several garbage collectors, as well as the design of the widely used region-based garbage collectors. I also performed a general discussion of the related work on analyzing region-based garbage collectors.

As discussed in this chapter, most garbage collectors generally have the similar design, especially for those region-based collectors which shares the common design of the heap structure and marking algorithm. This leads to an open question: In what extent does these region-based collectors relate to each other? To answer this question, in next chapter I will explore the differences and relationships among these collectors and implement them on JikesRVM as a series of progressive algorithmic improvements instead of build them individually.

Design and Implementation of the G1 Family of Garbage Collectors

As briefly discussde in the previous chapter, the four well-known region-based collectors, Garbage-First GC, Shenandoah GC, C4 GC and ZGC share many common design parts such as concurrent marking algorithm, region-based structure and allocation algorithm. However, there are also several differences among them, especially the different technique used by these collectors to improve GC performance. Hence their resulting GC performance can vary.

As a major contribution of thie thesis, this chapter explores the underlying relationships among these collectors which the original designers of these GCs was not aware of. Based on these discovered relationships, this chapter also discusses the implementations of several region-based GCs, as we called "The Garbage-First family of garbage collectors", as a series of improvements infered by the discovered relationships.

Section 3.1 describes the general and high-level design of the G1 family of garbage collectors and explores their underlying relationships. Section 3.2 describes the implementation details of a simple region-based GC. Section 3.3 describes the implementation of the simple region-based GC that uses linear scanning during evacuation. Section 3.5 describes the implementation of the improvement to the region-based GC that performs object marking concurrently. Section 3.5 describes the details of implementing the Garbage-First GC Section 3.6 describes the details of implementing the Shenandoah GC

3.1   General Design

The general design of the implementation includes the reconsideration of the underlying relationship among all the existing region-based garbage collectors, reorder them and implement them to reflect such relationships.

3.1.1   Analysis of the collectors and their relationships

The four well-known region-based garbage collectors, G1 GC, Shenandoah GC, C4 GC and ZGC share plenty of common design and algorithms. The major similarity of these collectors is that they are pauseless collectors originally designed for modern Java virtual machines and attempt to improve GC performance when targeting large heaps.

The similar design goal leads to the similar general structure of these collectors. As described in publications by Detlefs et al. [2004], Flood et al. [2016], Tene et al. [2011] and Liden and Karlsson [2018], their similarities include:

> **Region-based memory structure** which dividing the whole heap up into smaller regions and perform collection on a subset of regions for each GC (such subset is called the "collection set").
>
> **Bump pointer allocator** which uses a allocation algorithm that linearly allocates objects to the memory slice of a region and moves to another available region when the current region is filled.
>
> **Concurrent marking**. At the begining of each GC cycle, all of the collectors use the Snapshot-at-the-begining algorithm (firstly introduced by Yuasa [1990]) to mark all live objects in the heap in concurrent without stopping the mutators.
>
> **Perform heap compaction** after the marking phase copy live objects in the collection set to other regions before releases the memory in the collection set.

However, differences exist among these collectors. These garbage collectors use different technique to reduce the GC pause time and increase the mutator throughput. In terms of different collectors,

> Garbage-First GC uses a data structure called "remembered set" to record the cross region pointers during the execution of the Java program. Later during compaction phase, instead of walk over the whole object graph to update references as what other collectors (e.g. SemiSpace GC) did, G1 only needs to scan the remembered sets to update cross region pointers pointing to the collection set. Also G1 has a generational mode which eagerly collects newly allocated objects during nursery GCs.
>
> Shenandoah GC tries to perform the compaction phase in concurrent, by using the Brooks indirection barrier (which was firstly introduced by Brooks [1984]). Although Shenandoah GC still uses performs heap tracing during compaction phase, the GC pause time is extremely low and is not proportional to the heap size.
>
> ZGC can be simply regarded as a improvement to the Shenandoah GC with high mutator performance. ZGC uses the "colored pointers" as well

as a highly optimized read barrier to assist with the Brooks indirection barrier to further increase the mutator throughput.

C4 GC is very similar to ZGC and also performs concurrent compaction during each GC. The major difference to ZGC is that C4 GC is generational which is expected to have higher performance than ZGC.

In general, these four collectors are strongly related to each other in terms of the design and algorithm structure. In addition, each collector performs their own improvements to improve GC performance. In this way, at the algorithmic level these collectors tends to be a set of improvements to the naïve region-based collector. Hence they are more likely to form as a family of collectors instead of individual and unrelated collectors.

3.1.2   Implementation steps

As the first implementation which reflects such algorithmic hierarchy, members of the G1 family of garbage collectors are reorgainzed and reimplemented on JikesRVM by following different steps against the original implementation.

I started from implementing a simple region-based GC which divides the whole java heap up into fix sized regions and performs fully stop-the-world but parallel object marking and full-heap tracing based heap compaction/evacuation during each GC. Then I made progressive improvements based on this collector to further implement some G1 family of garbage collectors.

Since both Garbage-First and Shenandoah GC use heap linear scanning for evacuation, I implemented a linear-scan region-based GC by divided the heap evacuation phase of the simple region-based GC into two phases: The linear scan evacuation phase and the reference updating phase.

Then I changed the stop-the-world marking phase of the linear-scan region-based GC to the concurrent marking phase to implement a concurrent-marking region-based GC, by using the Snapshot at the Beginning algorithm (Yuasa [1990]).

By implementing the remembered-sets and remembered-set based evacuation based on the previous concurrent-marking region-based GC, I implemented the Garbage-First collector.

Also starting from the concurrent-marking region-based GC, by implementing the Brooks indirection pointers and the corresponding mutator barriors, I implemented the Shenandoah GC.

By performing such progressive improvements, I successfully implemented a series of G1 family of garbage collectors and share as much code and design as possible among the implementations of these collectors. Which enables the possibility for future algorithmic-level analysis on these garbage collectors.

3.2   Simple Region-based GC

This simple region-based GC is provided as a baseline for future implementation of the G1 family of collectors. It contains most basic structures of the Garbage-First family of garbage collectors such as region-based heap space and bump pointer allocation algorithm. By making progressive improvements on this region-based GC, I keep the implementation differences among the Garbage-First family of garbage collectors to a minimum.

3.2.1   Heap structure and allocation policy

This simple region-based GC and other GCs discussed later all uses the same implementation of the heap structure (which is implemented in `RegionSpace` and `Region` classes) and the same allocation policy (which is implemented in the `RegionAllocator` class). The `RegionSpace` divides the whole heap up into fixed 1 MB regions, which are 256 MMTk pages.

Figure 3.1 shows the code for allocation objects within the region space. To allocate objects, for each allocator, it firstly requests a region of memory (256 pages) from the page resource. Then it makes the allocation cursor points to the start of the region and bump increase this allocation cursor to allocate objects. Figure 3.1(a) shows the code for such allocation fast path. Most of the object allocation processes will only go to the fast path. However after a region is filled, to allocate a new object, the mutator enters a slow path which is shown in figure 3.1(b) that moves the bump pointer to another newly acquired region for future allocations.

MMTk reserves some extra pages for each region to record metadata. After the allocator filled a region, it records the end address of the region for region linear scanning which is later used in some collectors. Also, an off-heap bitmap is maintained in the metadata pages of the region to record the liveness data of objects allocated in this region.

3.2.2   Stop-the-world marking

Based on the high-level design of MMTk, after each allocation MMTk checks if a stop-the-world or concurrent GC is required for for the current selected GC plan. For this simple region-based GC, only stop-the-world collections are triggered.

The region-based GC initiates a collection cycle when the free heap size is less than the pre-defined `reversedPercent` (default is 10%). During each GC cycle, the collector starts by performing a full heap tracing to recursively mark all live objects. The marking algorithm considers the heap as a graph of objects and follows the idea of Breadth-first graph search which firstly scan and mark all the stack and global root objects and pushes them into an object queue. Then the collectors threads keeps popping objects from the object queue, collects all its object reference fields (which are child nodes of the current object node in the object graph), scan these fields and if they are not marked previously, push them back into the object queue. The marking

```
1   @Inline
2   public final Address alloc(int bytes, int align, int offset) {
3     /* establish how much we need */
4     Address start = alignAllocationNoFill(cursor, align, offset);
5     Address end = start.plus(bytes);
6     /* check whether we've exceeded the limit */
7     if (end.GT(limit)) {
8       return allocSlowInline(bytes, align, offset);
9     }
10    /* sufficient memory is available, so we can finish performing the allocation
          */
11    fillAlignmentGap(cursor, start);
12    cursor = end;
13    // Record the end cursor of this region
14    Region.setCursor(currentRegion, cursor);
15    return start;
16  }
```

(a) Region allocator - fast path

```
1   @Override
2   protected final Address allocSlowOnce(int bytes, int align, int offset) {
3     // Acquire a new region
4     Address ptr = space.getSpace(allocationKind);
5     this.currentRegion = ptr;
6
7     if (ptr.isZero()) {
8       return ptr; // failed allocation --- we will need to GC
9     }
10    /* we have been given a clean block */
11    cursor = ptr;
12    limit = ptr.plus(Region.BYTES_IN_REGION);
13    return alloc(bytes, align, offset);
14  }
```

(b) Region allocator - slow path

Figure 3.1: Region Allocator

```
1  @Inline
2  public ObjectReference traceMarkObject(TransitiveClosure trace, ObjectReference
       object) {
3    if (testAndMark(rtn)) {
4      Address region = Region.of(object);
5      // Atomically increase the live bytes of this reigon
6      Region.updateRegionAliveSize(region, object);
7      // Push into the object queue
8      trace.processNode(object);
9    }
10   return object;
11 }
```

Figure 3.2: Code for mark each object

process is done when the global object queue is drained, which means all object that are reachable from the root objects are marked after the marking phase.

In stead of using the gc byte in the object header for marking which is widely used in other GCs in MMTk, a bitmap for each region is maintained to record the liveness data for each object within this region. At the start of the marking phase, bitmaps of all regions are initialized to zero. As shown in figure 3.2, during visiting each object, the collector attempts to set the mark bit of the current object in the bitmap, and push the object into the object queue only if the attempt is succeed.

Although using such extra "live table" is not necessary for this region-based GC, this is a common design for Garbage-First and Shenandoah GC. So the gc byte in the object header was disabled at the very beginning to reduce the implementation difference among all collectors.

3.2.3   Collection set selection

As shown in figure 3.2, during the processing of each object in the heap, the collectors also atomically increase the live bytes for each region, starting from zero. After the full heap marking phase, the collector starts a collection set selection phase to construct a collection set of regions.

The collection set selection phase firstly takes a list of all allocated regions, sort them in ascending order by the live bytes of the region. Then the collector selects the regions with lowest live bytes. Also the collector should make sure the total live bytes in the collection set is not greater than the free memory size of the heap.

At the end of the collection set selection phase, the collector marks all regions in the collection set as "relocationRequired" for the future evacuation phase.

3.2.4   Stop-the-world evacuation

The evacuation phase is a fundamental part of the copying collectors. It tries to avoid heap fragmentation by forwarding the live objects which are sparsely located in the heap and compacting them together.

```
1  @Inline
2  public ObjectReference traceEvacuateObject(TraceLocal trace, ObjectReference
       object, int allocator) {
3    if (Region.relocationRequired(Region.of(object))) {
4      Word priorStatusWord = ForwardingWord.attemptToForward(object);
5      if (ForwardingWord.stateIsForwardedOrBeingForwarded(priorStatusWord)) {
6        // This object is forward by other threads
7        return ForwardingWord.spinAndGetForwardedObject(object, priorStatusWord);
8      } else {
9        // Forward this object
10        ObjectReference newObject = ForwardingWord.forwardObject(object, allocator
           );
11        trace.processNode(newObject);
12        return newObject;
13      }
14    } else {
15      if (testAndMark(object)) {
16        trace.processNode(object);
17      }
18      return object;
19    }
20  }
```

Figure 3.3: Code for evacuate each object

For this simple region-based GC, the evacuation phase is performed after the end of the collection set selection phase. This evacuation phase is aiming to copy/evacuate all live objects in the collection set to other regions.

Figure 3.3 shows the process of object evacuation. To evacuate objects, the collector performs another full heap tracing to remark all the objects, just like the marking phase discussed before. But in addition to mark the objects, the collector also copy the objects and atomically set the forwarding status in the object header if the object is in the relocation set. The same full heap tracing ensures all objects marked in the marking phase are also being scanned in the evacuation phase. Which means all live objects in the collection set are processed and evacuated properly.

At the end of the evacuation phase, after all objects are evacuated, the whole stop-the-world GC cycle is finished. The collector frees all the regions in the collection set and resumes the execution of all the mutators.

3.2.5   Evacuation correctness verification

Since the region-based evacuation is a fundamental component of G1 family of garbage collectors and can have several variants, it is necessary to verify the correctness of the evacuation process. I used an additional full heap tracing for verification. The full heap tracing is similar to the marking trace and is fully stop the world to ensure the object graph is never changed during the verification process. When visiting each object node during verification, the collector checks all its object reference fields and ensure they are either null or are pointing to the valid java objects. The collector

also checks no object node should be located in the from space (i.e. the collection set).

This full heap tracing verification process is optional and can be switched on and off for debugging purposes. At the end of the verification process, we can assert that the whole java heap is not broken and is in the correct status.

3.3   Linear Scan Evacuation

Since both Garbage-First and Shenandoah GC use heap linear scanning for evacuation, I implemented a linear-scan evacuation version of the simple region-based GC.

In the simple region-based GC, both object evacuation and reference updating are done together during the stop-the-world evacuation phase, using a single full heap tracing. For this linear-scan region-based GC, I divided the stop-the-world evacuation phase into two phases: The linear scan evacuation phase and the reference updating phase.

During the linear scan evacuation phase, the collector threads linear scans all the regions in the collection set and evacuate live objects in these regions. the collector does not fix or update any references during this phase.

During the reference updating phase the collector performs a full-heap tracing, just likes the evacuation phase of the simple region-based GC, but only fix and update pointers to ensure every pointer in the heap pointers to the correct copy of the object.

By separating the linear scan evacuation phase and the reference updating phase, the linear-scan region-based GC shows the basic collection processes of most G1 family of garbage collectors: `marking -> evacuation -> update references -> cleanup`. The modular design of this collectors enables the future redesign and rewrite of some specific phases, e.g. concurrent marking, remembered-set based evacuation or concurrent evacuation.

3.4   Concurrent Marking

Another improvement I did is to make the marking phase executes mostly in concurrent to reduce the mutator latency caused by the stop-the-world object marking. This concurrent-marking region-based GC is an improvement based on the previous linear-scan region-based GC.

The Snapshot at the Beginning (SATB) algorithm Yuasa [1990] is introduced to complete most of the marking work concurrently without pausing mutators. The SATB algorithm assumes an object is live if it was reachable (from the roots) at the start of the concurrent marking or if it was created during the concurrent marking Yuasa [1990].

The whole marking phase consists of two pauses, initial mark pause and final mark pause. During the initial mark pause, same as the stop-the-world marking phase, the collector clears the live bitmap for all regions, and then it scans and process all the root objects including stack objects and global objects. If the mutators allocate spaces too quick and the used ratio of the heap reaches the stop-the-world

GC threshold previously defined in the simple region-based GC, the collector then pauses all the mutators and switches to a stop-the-world marking phase to continue the marking process.

After the stop-the-world initial mark pause, the collector resumes the execution of all mutators and initiates the concurrent marking phase, which marks all the remaining objects in concurrent while the mutators are still in execution.

3.4.1   SATB write barriers

To maintain the invariant of "An object is live if it was live at the start of marking", the mutators in this concurrent-marking GC implemented a "deletion barrier", which are code fragments inserted into the java program, before every object field write, object field compare & swap and object array copy in the program. Figure 3.4 shows the object reference write barrier that is used to track object graph modifications during the concurrent marking. When a object reference field modification happens, since the old child object reference was reachable from the roots before this modification, it should be considered as a live object. So the SATB barrier traces and enqueues such old object reference fields to ensure they are live. Same as the stop-the-world marking process, the concurrent marking phase finishes after the global object queue is drained.

During the final mark pause, the collector releases and resets the marking buffer that is used for concurrent marking. Then it starts the collection set selection phase, just like the simple region-based GC.

This concurrent-marking region-based GC is an important improvement which enables the analysis of the pause time and mutator latency provided by the SATB concurrent marking algorithm on the region space. Although MMTk provides an implementation of SATB algorithm, but the old implementation is only performed on the mark-sweep space, not on the region space. By analyzing this concurrent-marking region-based GC, it is able to understand the detailed performance impact of the Garbage-First and Shenandoah GC due to the SATB algorithm.

3.5   Garbage-First GC

Garbage-First (G1) GC (Detlefs et al. [2004]) was originally designed by Oracle to replace the old Concurrent Mark and Sweep GC. G1 GC was designed to target large heap but has a reasonable and predictable GC pause time. To achieve a short pause time, G1 uses a data structure called remembered-set to perform partial heap scanning during the pointer updating phase, instead of the full-heap scanning. To make the GC time predictable, a pause time prediction model is involved in G1 to predict and choose the number regions in the collection set to meet a soft real-time pause goal. In addition, G1 has a generational mode which has some minor GCs that collects young (newly allocated) regions only.

```
1  @Override
2  protected void checkAndEnqueueReference(ObjectReference ref) {
3    if (!barrierActive || ref.isNull()) return;
4
5    if (Space.isInSpace(Regional.RS, ref)) Regional.regionSpace.traceMarkObject(
         remset, ref);
6    else if (Space.isInSpace(Regional.IMMORTAL, ref)) Regional.immortalSpace.
         traceObject(remset, ref);
7    else if (Space.isInSpace(Regional.LOS, ref)) Regional.loSpace.traceObject(
         remset, ref);
8    else if (Space.isInSpace(Regional.NON_MOVING, ref)) Regional.nonMovingSpace.
         traceObject(remset, ref);
9    else if (Space.isInSpace(Regional.SMALL_CODE, ref)) Regional.smallCodeSpace.
         traceObject(remset, ref);
10   else if (Space.isInSpace(Regional.LARGE_CODE, ref)) Regional.largeCodeSpace.
         traceObject(remset, ref);
11 }
```

(a) The SATB Barrier which enqueues objects into a SATB buffer

```
1  @Inline
2  public void objectReferenceWrite(ObjectReference src, Address slot,
       ObjectReference tgt, Word metaDataA, Word metaDataB, int mode) {
3    if (barrierActive) checkAndEnqueueReference(slot.loadObjectReference());
4    VM.barriers.objectReferenceWrite(src, tgt, metaDataA, metaDataB, mode);
5  }
6
7  @Inline
8  public boolean objectReferenceTryCompareAndSwap(ObjectReference src, Address
       slot, ObjectReference old, ObjectReference tgt, Word metaDataA, Word
       metaDataB, int mode) {
9    boolean result = VM.barriers.objectReferenceTryCompareAndSwap(src, old, tgt,
         metaDataA, metaDataB, mode);
10   if (barrierActive) checkAndEnqueueReference(old);
11   return result;
12 }
13
14 @Inline
15 public boolean objectReferenceBulkCopy(ObjectReference src, Offset srcOffset,
       ObjectReference dst, Offset dstOffset, int bytes) {
16   Address cursor = dst.toAddress().plus(dstOffset);
17   Address limit = cursor.plus(bytes);
18   while (cursor.LT(limit)) {
19     ObjectReference ref = cursor.loadObjectReference();
20     if (barrierActive) checkAndEnqueueReference(ref);
21     cursor = cursor.plus(BYTES_IN_ADDRESS);
22   }
23   return false;
24 }
```

(b) SATB mutator barriers

**Figure 3.4:** Snapshot at the Beginning Barriers

Figure 3.5: Remembered-Set Structure

During each GC cycle of G1 GC, there are 5 major phases. The first three phases are the concurrent marking phase, the collection set selection phase and the linear scan evacuation phase, same as the concurrent-marking region-based GC. The fourth phase is the remembered-set based pointer updating phase. And the last phase is the cleanup phase which frees the regions in the collection set.

The construction of the Garbage-First GC on JikesRVM is based on the concurrent-marking region-based GC. It involves three improvements to the concurrent region-based collector discussed in : remembered-set based pointer updating, pause time predictor and generational collection. The remembered-set based pointer updating largely reduces the pause time cause by the reference updating phase by avoiding full heap tracing when reference updating, using remembered sets. The pause time predictor uses a statistical prediction model to make the pause time of each GC more predictable and not exceeds a soft pause time goal at most of the time. The generational mode enables G1 to collect garbage more efficiently, under the help of the generation assumption of objects, which under most cases further reduces the average pause time of GCs

3.5.1   Remembered-set

Under the G1 collection policy, the region space further divides regions into fixed 256 B cards for constructing remembered sets.

Figure 3.5 show the general structure of a remembered-set. The remembered-set is a data structure for each region that remembers cross region pointers in other regions that pointing to objects in the current region. A remembered-set for a region is implemented as a list of PerRegionTables. A PerRegionTable in the remembered-set is a bitmap corresponds to a region in the heap. The bitmap records cards in the corresponding region that contains pointers pointing to the region that the current remembered-set corresponds to. Each bit in the bitmap corresponds to one card.

The remembered-sets is maintained by the collector and is always contains cards in other regions that has pointers to this region. The remembered-sets should be updated at every object field modification bu mutators and after the evacuation of

```
1   @Inline
2   void markAndEnqueueCard(Address card) {
3     if (CardTable.attemptToMarkCard(card, true)) {
4       remSetLogBuffer().plus(remSetLogBufferCursor << Constants.
          LOG_BYTES_IN_ADDRESS).store(card);
5       remSetLogBufferCursor += 1;
6       if (remSetLogBufferCursor >= REMSET_LOG_BUFFER_SIZE) {
7         enqueueCurrentRSBuffer(true);
8       }
9     }
10  }
11
12  @Inline
13  void checkCrossRegionPointer(ObjectReference src, Address slot, ObjectReference
        ref) {
14    Word x = VM.objectModel.objectStartRef(src).toWord();
15    Word y = VM.objectModel.objectStartRef(ref).toWord();
16    Word tmp = x.xor(y).rshl(Region.LOG_BYTES_IN_REGION);
17    if (!tmp.isZero() && Space.isInSpace(G1.G1, ref)) {
18      Address card = Region.Card.of(src);
19      markAndEnqueueCard(card);
20    }
21  }
```

Figure 3.6: Remembered-set Barrier

each object in the collection set to meet the correctness of the remembered-sets.

3.5.2 Concurrent remset refinements

As shown in figure 3.6, in order to maintain the structure of remembered-sets, a new barrier called "remembered-set barrier" was involved for each object field modification action in the java program. For each object field modification `obj.x = y`, the remembered-set barrier checks whether the pointer `y` is a cross region pointer (i.e. not pointing to the region that contains `obj`). If the check succeed, the remembered-set barrier enters a slow path which marks the card containing `obj` and push this card into a local `dirtyCardBuffer`. Since the minimum memory allocation unit in the MMTk metadata space is one page (4 KB), the size of the local `dirtyCardBuffer` is 1024 cards, instead of 256 cards in the original G1 implementation in OpenJDK Detlefs et al. [2004].

When the local `dirtyCardBuffer` is full, the remembered-set barrier pushes the local `dirtyCardBuffer` to the global filled RS buffer set. And when the size of the global filled RS buffer set reaches a threshold of 5 `dirtyCardBuffers`, a concurrent remset refinement is triggered. A separate concurrent remset refinement thread was started to process each `dirtyCardBuffer` in the global filled RS buffer set. The refinement thread scans each card for each `dirtyCardBuffer`. If the card is marked, clear its marking data, linear scan it for cross region pointers and updates the corresponding remembered-set for each cross region pointer.

The card table

To perform card marking and card linear scanning during concurrent remset refinements, A card table should be used to record the marking data for each card as well as the offset for the first and the last object of each card. The card table consists of three parts. a) A bitmap of all cards in the java heap to record the marking data of the cards. b) A byte array (i.e. a byte map) to record the offset of the first object for each card. c) Another byte array to record the offset of the last object for each card.

Hot cards optimization

During concurrent remset refinements, some cards may be enqueued and scanned multiple times which brings extra computation costs. To avoid such redundant scanning, a hotness value are assigned for each card. Every scan on a card increases its hotness value by 1. When the hotness for a card exceeds a hotness threshold (default is 4), this card is considered as a "hot card". Hot cards are pushed into a hot cards queue and the processing of all the hot cards card are delayed until the start of the evacuation phase.

3.5.3 Evacuation phase

G1 uses linear scan evacuation, just like the previous discussed collectors. During the evacuation phase, the collector linear scans each region in the collection set and copy live objects to the to-regions.

3.5.4 Remembered set based pointer updating

By performing concurrent remset refinements, the structure of all remembered-sets are always maintained correctly. Which enables the partial evacuation based on remembered sets without a full heap tracing.

After the collection set selection phase, the collector firstly performs a linear scan over regions in the collection set to evacuate all love objects in the collection set. Then the collector starts a references updating phase which considers the root set as all root objects plus objects in the cards that is recorded in the remembered-sets. During references updating phase, instead of performing a full heap tracing to fix and update all the references in the heap, the collector only scans root objects and cards in remembered-sets of regions in the collection-set to update references, since all pointers that needs to be updated are either root pointers or are remembered in the remembered-sets.

At the end of the pointer updating phase, same as the full-heap tracing based pointer updating, the collector frees all the regions in the collection set.

By the remembered set based pointer updating, the G1 collector has the ability to collect any subset of regions in the heap, without scanning the whole heap, as long as the collection set size is not greater than the free memory size in the heap. In this way the GC pause time due to object evacuation can be largely shortened.

3.5.5 Pause time predictor

Due to the ability of collecting any subset of regions in the heap, G1 can further make the pause time of each GC predictable by choosing the number of regions to in the collection set.

The total kinds of works during the stop-the-word evacuation is fixed: dirty cards refinements, object evacuation and linear scan cards for updating references. This brings us the ability to model the pause time for each GC, in terms of the remembered-set size and live bytes of each region in the collection set.

Our implementation reuses the pause time prediction model that was proposed in the original G1 paper (Detlefs et al. [2004]).

$$T_{CS} = T_{fixed} + T_{card} * N_{dirtyCard} + \sum_{r \in CS} (T_{rs\ card} * rsSize(r) + T_{copy} * liveBytes(r))$$

Where
T_{CS} is the total predicted pause time
T_{fixed} is the time of all extra works
T_{card} is the time of linear scanning a card for remembered set refinements
$N_{dirtyCard}$ is the number of dirty cards that has to be processed before evacuation
$T_{rs\ card}$ is the time of linear scanning a card in the remembered-set for evacuation
$rsSize(r)$ is the number of cards in the remembered-set of region r
T_{copy} is the time for evacuating a byte
$liveBytes(r)$ is the total live bytes for evacuation

By using this pause time prediction model, during each collection set selection phase, the collector can choose the number of regions in the collection set to meet a user-defined pause time goal. This mechanism makes the pause time more predictable and being controlled around a reasonable value.

3.5.6 Generational G1

The original design of G1 GC comes with a generational mode which collects newly allocated regions only during young GCs (Detlefs et al. [2004]). The generational collection is based on an assumption that the newly allocated objects has higher chances to become garbage than those objects that is survived after several GCs.

Based on this assumption on the age of objects, the Generational G1 collector divides the regions into three generations: Eden, Survivor and Old generation, as shown in figure 3.7. Eden regions contain objects that is newly allocated since last GC. Survivor regions contain objects that is survived from last GC. And Old regions contain objects that is survived after several GCs. During the allocation process, the newly allocated regions are marked as Eden regions. When the ratio of the number of Eden regions exceeds a `newSizeRatio` threshold, a young GC is triggered which only collects all Eden and Survivor regions. During young GCs, live objects in

Figure 3.7: Generational G1 Heap Structure

Eden regions are evacuated to Survivor regions and objects in Survivor regions are evacuated to Old regions. Objects evacuated to Survivor regions will still be included into the collection set during the next young GC.

Young GC

Young GCs are just simple nursery GCs that only evacuate objects and not marking them. There is no marking phase during a young GC. Instead of determine the liveness of objects by marking, during young GCs the collector considers all objects that are not in the to-space (i.e. the collection set) as live. The collector simply starts from the the root objects and the remembered-set to recursively evacuate live objects out of the collection-set.

Mixed GC

When the allocate memory in the heap exceeds some threshold, the generational G1 will initiate a concurrent marking phase for a mixed GC just likes the non-generational G1. When the to-space is exhausted during mixed GCs, G1 switches to a fully stop-the-world full GC.

Pause time predictor for young GCs

To meet the soft pause time goal during young GCs, the collector update the value of `newSizeRatio` at the end of every young GC to find a more appropriate Eden size ratio so that the collector can performs better in meeting the pause time goal during future young GCs.

3.6   Shenandoah GC

Shenandoah GC is a experimental garbage collector and is originally implemented on OpenJDK. Shenandoah GC is designed to reduce GC pause times for large heaps and tries to make the GC pause time not proportional to the heap size.

When the heap occupancy reaches a threshold ratio (default is 20%), the Shenandoah GC triggers a concurrent GC cycle, starting from a concurrent marking phase. The concurrent marking phase also uses the Snapshot at the Beginning algorithm,

which is similar to the G1 GC and the concurrent-marking region-based GC. After the concurrent marking phase is a concurrent remembered-set selection phase which is the same as the stop-the-world remembered-set selection phase but runs in concurrent, without pausing the mutators. The third phase is the concurrent evacuation phase. Shenandoah GC uses the Brooks-style indirection pointer (Flood et al. [2016]) to evacuate objects in to-space in concurrent. The fourth phase is the concurrent reference updating phase, which also performs a full heap tracing, like the stop-the-world reference updating phase, but runs in concurrent, under the help of the Brooks-style indirection pointers.

By making the marking, evacuation and pointer updating phases run in concurrent, Shenandoah GC performs most of the heap scanning works in concurrent, without pausing mutators. In this way Shenandoah GC can have very low pause times and the pause time is not proportional to the heap size.

3.6.1   Brooks indirection barrier

The Brooks-style indirection pointer is a pointer stored in the java object header to record the forwarding status of a java object (Flood et al. [2016]). This requires the Shenandoah collector in JikesRVM to reserve an extra GC word for each java object. During object allocation, after a new java object is allocated, the indirection pointer of this java object is initialized to point to the object itself.

During the concurrent evacuation phase, after an object is forwarded, the collector atomically update the indirection pointer of the old copy to point to the new copy. Mutators should always perform modification actions on the new copy of the object, by following the indirection pointer in the object header.

By using the Brooks-style indirection pointers, the mutators can read and modify object (by following the indirection pointers) while the collector can also evacuate these objects in concurrent.

Read Barriers

The read operations of every object field `obj.x`, including the object reference fields and the primitive fields, should go through the object's indirection pointer. Figure 3.8 shows the read barrier that is inserted before every object field read instruction. The barrier always unconditionally extract the indirection pointer from the object header, and reads the object field value from the object that the indirection pointer points to.

If the object is not in the collection set or is in the collection set but is not forwarded, reading data from its original copy is safe since there is no other copies of the object at the time the object field read happens. If the object is forwarded, its indirection pointer is pointed to the new copy of the object. Then by following the indirection pointer, the mutator can always read field data from the object's new copy.

```
1   // class ShenandoahMutator
2
3   @Inline
4   public ObjectReference getForwardingPointer(ObjectReference object) {
5      return ForwardingWord.extractForwardingPointer(object);
6   }
7
8   @Inline
9   public ObjectReference objectReferenceRead(ObjectReference src, Address slot,
          Word metaDataA, Word metaDataB, int mode) {
10     return VM.barriers.objectReferenceRead(getForwardingPointer(src), metaDataA,
             metaDataB, mode);
11  }
12
13  // class ForwardingWord
14
15  @Inline
16  public static ObjectReference extractForwardingPointer(ObjectReference oldObject
          ) {
17     return oldObject.toAddress().loadWord(FORWARDING_POINTER_OFFSET).and(
             FORWARDING_MASK.not()).toAddress().toObjectReference();
18  }
```

Figure 3.8: Brooks read barrier

Write Barriers

However, since the mutator should always meet the "write in the to space" invariant, it is not safe for mutators to write objects just by following the indirection pointers. If a mutator is trying to updating a object while a collector thread is also copying this object, the write operation will only performs on the old copy and the new copy may still contains the old data, because the indirection pointer is still pointing to the old copy when the evacuation of this object is in progress.

To resolve this problem, the write barrier used in Shenandoah GC is different from the read barrier. As shown in figure 3.9, before the mutator performs the object field write operation on an object that is in the collection set, the mutator firstly checks if the indirection pointer of the object is marked as FORWARDED or BEING_FORWARDED. If the object is forwarded, then the mutator follows the indirection pointer to perform the write operation on the new copy. If the object is marked as being forwarded, the mutator awaits until the object is forwarded to get the correct indirection pointer to the new copy. If the object is marked as not forwarded, the mutator should take the responsibility to forward this object. Under such situation, the mutator copies this object and updates the indirection pointer, just like the forwarding process of the collector threads.

After the mutator takes the responsibility of forwarding unforwarded from-space objects, the mutator now meets the "write in the to space" invariant.

```
1  @Inline
2  public ObjectReference getForwardingPointerOnWrite(ObjectReference object) {
3    if (brooksWriteBarrierActive) {
4      if (isInCollectionSet(object)) {
5        ObjectReference newObject;
6        Word priorStatusWord = ForwardingWord.attemptToForward(object);
7        if (ForwardingWord.stateIsForwardedOrBeingForwarded(priorStatusWord)) {
8          // The object is forwarded by other threads
9          newObject = ForwardingWord.spinAndGetForwardedObject(object,
               priorStatusWord);
10       } else {
11         // Forward this object before write
12         newObject = ForwardingWord.forwardObjectWithinMutatorContext(object,
               ALLOC_RS);
13       }
14       return newObject;
15     } else {
16       return object;
17     }
18   } else {
19     return getForwardingPointer(object);
20   }
21 }
22
23 @Inline
24 public void objectReferenceWrite(ObjectReference src, Address slot,
       ObjectReference tgt, Word metaDataA, Word metaDataB, int mode) {
25   ObjectReference newSrc = getForwardingPointerOnWrite(src);
26   VM.barriers.objectReferenceWrite(newSrc, tgt, metaDataA, metaDataB, mode);
27 }
```

Figure 3.9: Brooks write barrier

```
1   @Override
2   @Inline
3   public void processEdge(ObjectReference source, Address slot) {
4     ObjectReference oldObject, newObject;
5     do {
6       oldObject = slot.prepareObjectReference();
7       newObject = traceObject(oldObject);
8       if (oldObject.toAddress().EQ(newObject.toAddress())) return;
9     } while (!slot.attempt(oldObject, newObject));
10  }
```

Figure 3.10: Concurrent update references

3.6.2   Concurrent evacuation

The evacuation phase is done mostly in concurrent. Before concurrent evacuation, the collector firstly collects and evacuate all root objects stop-the-world. Then during concurrent evacuation, the collector linear scans each region in the collection set to evacuate all live objects and atomically set their indirection pointers.

JikesRVM has its own implementation of object forwarding functions, but it stores the forwarding pointer into the status word in the object header. But by using the Brooks-style indirection pointer, all objects should have a valid indirection pointer in the header, which will override other status information, e.g. lock and dynamic hash status. So I use a new implementation of object forwarding functions which increases the java header by one word and stores the indirection pointer in this extra word.

Monitor objects

Monitor objects is an exception that is handled specially in the JikesRVM implementation of the Shenandoah GC. In JikesRVM, the jvm `monitorenter` and `monitorexit` instruction does not triggers read barriers when locking objects, which can cause inconsistencies of the lock status of the object. I modified the JikesRVM's monitor lock and unlock code to make them trigger the read barriers.

3.6.3   Concurrent updating references

Shenandoah GC also performs the update references phase in concurrent. The implementation is similar to the concurrent marking phase, but in addition to concurrently marks objects again, the collector also updates the object reference fields for each objects and make them points the correct object.

Since the java program is running during the concurrent reference updating phase, unconditionally update object reference fields may cause racing problems. So instead of unconditionally updating pointers which was used in the previous stop-the-world GCs, I modified the procedure to perform atomic pointer updating, as shown in figure 3.10, to avoid racing problems.

```
1  @Inline
2  public boolean objectReferenceCompare(ObjectReference lhs, ObjectReference rhs)
       {
3    if (lhs.toAddress().NE(rhs.toAddress()) && getForwardingPointer(lhs).toAddress
         ().NE(getForwardingPointer(rhs).toAddress())) {
4      return false;
5    } else {
6      return true;
7    }
8  }
```

Figure 3.11: Object Comparison Barrier

Object comparisons

Since the mutator can either load a old or a new object reference of a same object from some object fields, simply comparing addresses of objects for the `if_acmpeq` and `if_acmpne` instruction can cause false negatives. I implemented a new object comparison barrier, as shown in figure 3.11, to compare object references. In stead of simply comparing the addresses of the objects, the object comparison barrier also compares the indirection pointers of the objects.

The comparison of the original object addresses is still required. If we only compare the indirection pointers $a' == b'$, the GC can update the indirection pointer b' to a new address after the barrier extracted a' and before extracting b', which can still cause false negatives if a and b is the same object.

The object comparison barrier is implemented for both JikesRVM's baseline and optimizing compiler.

Object compare and swaps

When the java mutator is performing object reference compare and swap operations, the collector threads can also updating the pointers contained by the slot that the CAS is operating on. Under such situation the CAS may failed even when there is only one mutator thread updating the slot, since the old value hold by the object field slot is updated by the collector in concurrent.

Figure 3.12 shows the object compare and swap barrier used in Shenandoah GC. Instead of simply compare and swap on the old object, after the CAS on the old object failed, another CAS operation is performed where the old value is the indirection pointer of the old object.

3.6.4   Concurrent cleanup

The last phase of Shenandoah GC is the concurrent cleanup phase, during which the collector concurrently visits each region in the collection set, clear their metadata and release these regions. This phase runs in concurrent without pausing mutators.

```
1  @Inline
2  public boolean objectReferenceTryCompareAndSwap(ObjectReference src, Address
       slot, ObjectReference old, ObjectReference tgt, Word metaDataA, Word
       metaDataB, int mode) {
3    boolean result = VM.barriers.objectReferenceTryCompareAndSwap(src, old, tgt,
         metaDataA, metaDataB, mode);
4    if (!result) {
5      result = VM.barriers.objectReferenceTryCompareAndSwap(src,
           getForwardingPointer(old), tgt, metaDataA, metaDataB, mode);
6    }
7    return result;
8  }
```

Figure 3.12: Object Reference Compare and Swap Barrier

3.7 Summary

This chapter explores and discusses relationship of the G1 family of garbage collectors. Then each step of improvement I performed to the collector and the involved GC algorithm is also explained and discussed, as well as a little bit further discussion of the pros & cons for each structural component at an algorithmic level

However, the pros and cons of several algorithms discussed in this chapter are all discussed based on theorical reasoning. The real-world GC performance impact caused by these improvements still remains unexplored. So to explore the real-world GC performance, in next chapter I will perform several performance measurements of these collectors, on a set of banchmark suite that represents the real-world Java program.

Performance Evaluation

Answering the question of the GC performance under a real-world setting, this chapter discusses the performance evaluation I undertaken for analyzing the Garbage-First family of garbage collectors. This chapter will firstly describe the software and hardware platform I used for benchmarking, Then the detailed evaluation process and resutls of both pause time and barrier latency will be presented and discussed.

Section 4.1 and 4.2 roughly introduce the experimental platform and software involved for the performance measurements. Section 4.3 discusses the generall evaluation method used to evaluate the performance of all the interested metrics. Section 4.4, 4.5 and 4.6 discusses the detailed measurement steps involved to evaluate the GC pause time, barrier overhead and remembered set size respectivelt, as well as figures of the results and the discussion based on the results.

4.1   The Dacapo Benchmark

The Dacapo Benchmark Suite is a tool for Java benchmarking and contains a set of open-sourced real world programs with a high memory loads.

The Dacapo Benchmark is frequently used during the development of the G1 family of garbage collectors in chapter 3, as a validation program to verify the correctness of the collectors under a real world setting.

I also performed pause time evacuations and barrier latency evaluation on all of the following Dacapo Benchmark suites. The benchmarking suites I used for evaluation includes (Blackburn et al. [2006]):

- **luindex** Uses lucene to indexes a set of documents; the works of Shakespeare and the King James Bible

- **bloat** Performs a number of optimizations and analysis on Java bytecode files

- **hsqldb** Executes a JDBCbench-like in-memory benchmark, executing a number of transactions against a model of a banking application

- **lusearch** Uses lucene to do a text search of keywords over a corpus of data comprising the works of Shakespeare and the King James Bible

Machine	bobcat	rat	fisher	ox
CPU Info	AMD FX-8320	Intel i7-4770	Intel i7-6700k	Intel Xeon Gold 5118
Clock	3.5GHz	3.4GHz	4GHz	2.3GHz
Processors	4/8	4/8	4/8	4/48/96
L2 Cache	4MB	4MB	1MB	48MB
RAM	8GB	8GB	16GB	512GB

Table 4.1: Machines used for development and evaluation.

- **pmd** Analyzes a set of Java classes for a range of source code problems

- **xalan** Transforms XML documents into HTML

- **avrora** Simulates a number of programs run on a grid of AVR microcontrollers

- **sunflow** Renders a set of images using ray tracing

By performing evaluation on a wide range of benchmarking suites which represents different class of real world programs, it is more possible to understand the pros and cons of the G1 family of collectors under a real world setting.

4.2   Hardware Platform

During the implementation of all the G1 family of garbage collectors in chapter 3, a list of machine with a large varities on CPU types, clock, number of processors and the size of cache and memory were involved, as shown in Table 4.1. By executing all the benchmarking suite of the Dacapo Benchmark on these different machines, and thanks to the benchmarking suites of the Dacapo Benchmark which reflect that different categories of programs in the real world, we have the ability to statistically verify the correctness of the prevously implemented garbage collectors (in chapter 3) and make sure they performs as intended in a real world setting.

For further performance evaluation, the "fisher" machine was used for the final benchmarking.

4.3   General Measurement Methodology

I perform both GC pause time and barrier overhead measurements on the optimized build of the JikesRVM, with 6 different GC configurations for the 6 GCs to be measured respecively. To collect measurement results, I execute all 10 dacapo benchmark suites discussed in section 4.1 on all 6 collectors, with respect to 4 different heap sizes, 637 MB, 939 MB, 1414 MB and 1971 MB respecively. Each benchmark suite is executed 10 times for each configuration of ($GC, HeapSize$) to collect more precise results and avoids the error due to some unexpected environmental fluctuations.

4.3.1   Reducing non-determinisms

The adaptive compiler can have non-deterministic behaviours when performing dynamic compilations and optimizations to the executing Java program. In order to minimize such non-deterministic behaviours and makes the program executes more faster, I performed the measurement methodology called "warmup replay", which was firstly introduced by Yang et al. [2012], as a replacement of the "pseudoadaptive approach" introduced by Blackburn and Hosking [2004].

This methodology performs a execution of 10 iterations for each benchmark suite to collect runtime execution informations before the measurement, to assist with more optimized compilation. Then during measurements, after the first iteration of warmups, JikesRVM compiles all methods by using the advice information generated previously to avoid any re-compiling behaviours buring the following measurement iteration.

I ran all of the 9 benchmark suites discussed in section 4.1 on each collector for 10 times, with 2 times of warmup execution before the timing iteration. JikesRVM performs warmup-replay compilation at the end of the first warmup iteration.

4.4   Pause Time Evaluation

This section describes the steps took for pause time evaluation as well as all the evaluation results and discussions.

4.4.1   Mutator latency timer

In order to perform more careful analysis on the mutator pause times, instead of simply calculating the time starting from the first stop-the-world phase to the last stop-the-world phase during each GC cycle, I implemented a mutator latency timer to perform more precise calculation of mutator pause times.

The mutator latency timer contains a static "three dimemsional" array:

```
static long[] LOGS = long[THREAD ID * EVENT ID * TIMESTAMPS]
```

This array is statically allocated within the VM Space to record the timestamp (in nanoseconds) of each event, for each mutator. The first dimension is the thread id of all mutators. The second dimension is the event id. The third dimemsion TIMESTAMPS is the max number of logs of one event and is currently set to 1024. Particularly under current context, to measure the pause time for each mutator, two events, MUTATOR_PAUSE and MUTATOR_RESUME are defined to record the time when a mutator thread starts waiting for gc complete and the time when a mutator gets resumed for execution.

In JikesRVM, a mutator thread checks for GC requests and starts waiting if necessary every time it reaches a yieldpoint, which may trigger a MUTATOR_PAUSE event if it should be paused. After a stop-the-world cycle is finished, before continuing for further execution, it triggers a MUTATOR_RESUME event.
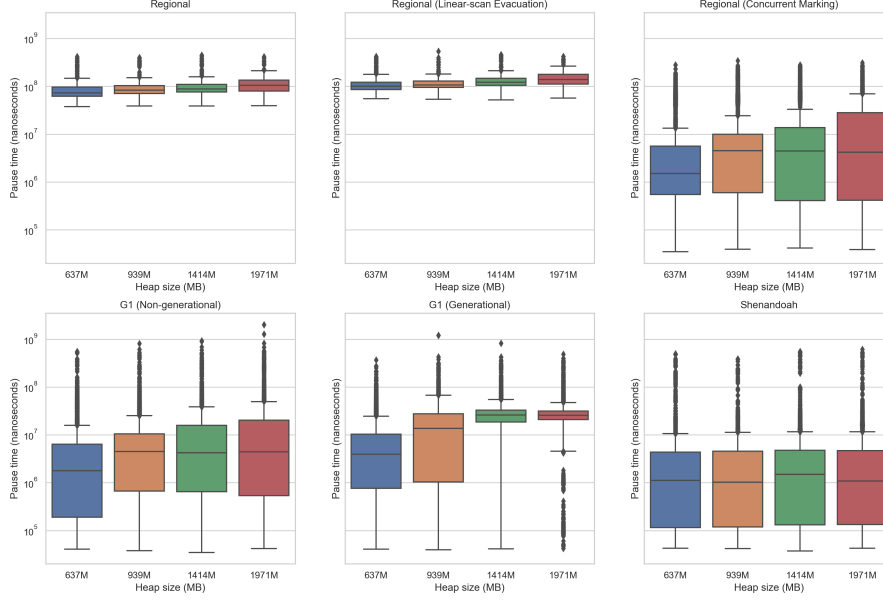
Figure 4.1: Pause times of 6 collectors

At the end of the benchmark execution, the mutator latency timer will dump all the data in the `LOGS` array to the print buffer for further data analysis.

As an output of the analysis of the overhead data, I report the minimun, 25%, 50%, 75% and maximum mutator pause time for each GC, each benchmark suite and each heap size.

4.4.2 MMTk harness callbacks

During each iteration of benchmarking, the Dacapo benchmark has several warm-up executions which will run the benchmark suite a few times to warm up the cache and JVM. Then dacapo will starts the actual benchmarking run. I use a probe called `MMTKCallback` which will call the `org.mmtk.plan.Plan.harnessBegin` method before the final benchmarking execution and call the `org.mmtk.plan.Plan.harnessEnd` after the final benchmarking execution. Based on this, the two callbacks `harnessBegin` and `harnessEnd` are used to calculate the inform the mutator latency timer to start recording logs or dump all logs.

4.4.3 Results

Figure 4.1 shows the results of the GC latency time evaluated on all 6 garbage collectors discussed in section 3. Each subfigure shows the overall GC latency time for each garbage collector. Specifically, the minimum, 25% percentile, medium, 75% percentile

Heap Size = 637 MB				
Pause time (ms)	Mean	Min	Mid	95%	Max
Regional	85.74	37.92	73.26	138.55	416.36
Regional Linear-scan evacuation	110.06	55.28	101.05	160.78	424.84
Regional Concurrent marking	11.36	0.04	1.51	67.06	279.70
Garbage-First Non-generational	9.51	0.04	1.79	46.69	548.71
Garbage-First Generational	10.36	0.04	3.92	44.21	364.37
Shenandoah	4.12	0.04	1.11	9.36	488.65

Heap Size = 939 MB				
Pause time (ms)	Mean	Min	Mid	95%	Max
Regional	94.40	39.07	82.65	137.98	401.14
Regional Linear-scan evacuation	117.74	53.42	107.33	171.03	545.52
Regional Concurrent marking	18.74	0.04	4.50	74.76	345.44
Garbage-First Non-generational	11.52	0.04	4.46	45.52	803.90
Garbage-First Generational	18.45	0.04	13.63	50.73	1198.33
Shenandoah	3.57	0.04	1.02	10.35	385.64

Heap Size = 1414 MB				
Pause time (ms)	Mean	Min	Mid	95%	Max
Regional	100.92	38.83	88.16	187.50	449.50
Regional Linear-scan evacuation	131.52	51.83	121.82	197.04	463.21
Regional Concurrent marking	21.18	0.04	4.43	82.69	281.24
Garbage-First Non-generational	13.40	0.03	4.20	43.27	926.35
Garbage-First Generational	27.82	0.04	25.98	54.36	820.84
Shenandoah	5.28	0.04	1.50	13.63	549.97

Heap Size = 1971 MB				
Pause time (ms)	Mean	Min	Mid	95%	Max
Regional	114.97	39.53	105.31	215.11	415.50
Regional Linear-scan evacuation	150.84	56.66	138.37	234.80	420.24
Regional Concurrent marking	23.98	0.04	4.19	93.37	310.79
Garbage-First Non-generational	15.52	0.04	4.38	46.07	2010.66
Garbage-First Generational	29.20	0.04	25.52	52.58	487.91
Shenandoah	5.55	0.04	1.08	18.06	613.34

Table 4.2: Results of the GC pause time

Overheads (ns)	Garbage-First Non-generational	Garbage-First Generational
637M	10.26% ±26.26%	5.72% ±15.21%
939M	5.71% ±14.52%	2.24% ± 7.74%
1414M	4.75% ±10.95%	2.38% ± 6.76%
1971M	9.41% ±46.04%	2.21% ± 6.94%
Overall	9.56% ± 7.09%	8.36% ± 6.99%

Table 4.3: Ratio of full gc

and maximum GC latency time nanoseconds were reported for each collector and each heap size.

To present the results more clearly, the benchmarking results ar presented as a set of box plots and with log-scaled pause time y in nanoseconds as y-axis.

4.4.4 Discussion

As the most simple form of the G1 family of GC, the Regional collector, performs fully stop-the-world GC during each GC cycle. Works for each GC cycle includes marking all live objects and walk over the object graph to evacuate objects in a set of selected memory regions. Totally 2 full heap tracing are performed during each GC cycle, which makes the GC pause time longer. The average pause time for this Regional GC generally ranges from 104 to 147 milliseconds and inceases as the heap size inceases.

By performing linear scan based evacuation, the linear-scan evacuation version of the Regional GC has to scan the memory regions in the collection set twice, one for exacuating live objects and one for updating references. In this way the collector generally has longer pause times, which ranges from 139 to 198 milliseconds in average and and inceases as the heap size inceases. The percentage of GC pause time increase are 33.7%, 30.5%, 26.9%, 34.8% for 4 heap sizes respecively, in average 31.5%. Although performing linear scan evacuation can largely increases the GC pause time, this indenpendent phase is a important part for G1 GC and the Shenandoah GC.

After performing concurrent marking in additional to the linear scan based evacuation, the resulting ConcRegional collector split the pause for marking into several smaller pauses and performs most of the marking work in concurrent without stopping the mutators. This makes the total pause time for a GC smaller and significantly reduces the average pause time by 88.5%.

The non-generational G1 reduces the collection size to meet a pause time goal of 100 ms. Based on the benchmarking results, at least 75% of the pauses are less than the pre-defined pause time goal. On the other hand, G1 uses remembered sets to update references instead of performing full heap tracing. To small heaps this has almost no benefits, even reduces the average pause time by 24.9% on 637 MB heaps. But for heap sizes of 39 MB, 1414 MB and 1971 MB, this partial heap scanning technique

reduces average pause time by 8.8%, 7.4% and 10.7% respecively. So the remembered sets based references updating has more benefits on larger heaps. However, the full GC becomes more expensive because of the large work required to update all the remembered sets in the heap, which usually results in a pause time ranges from 0.5 s to 1.5 s.

By using the generational G1 GC, young/nursery GCs are usually triggered several times before a major GC happens. Also nursery GCs are fully stop-the-world and always tries to collect as much nursery regions as possible, as long as the pause time does not exceeds the pre-defined pause time goal. This results in the incease of GC pause times. However, generational collection can largely reduces the probability of G1 GC falling back to full GCs. As shown in 4.3, I measured the overall full GC ratio for both non-generational and generational G1 GC on four different heap sizes, with 95% confidence interval reported as well. Based on the results, with generational mode, G1 reduces the probability of falling back to full GC by 12.5%.

The Shenandoah GC performs marking, evacuation and reference updating in concurrent, this significantly reduces the pause time by 72.6%, compared to the concurrnet marking version of the regional collector. Based on the benchmarking results at least 75% of the GC pauses do not exceeds 10 milliseconds. However, full GCs can still results in pauses of around 500 to 600 milliseconds, which is longer than the concurrent-marking regional collector due to the Brooks barrier involved during evacuation, which performance will be discussed in section 4.5

4.5 Barrier Latency Evaluation

This section describes the steps took for barrier overhead evaluation as well as all the evaluation results and discussions.

4.5.1 Methodology

The mutator barrier percentage overhead is modelled as follows:

$$\text{Overhead} = \frac{|\text{Mutator time with barrier} - \text{Mutator time without barrier}|}{\text{Mutator time without barrier}} * 100\%$$

The mutator execution time is calculated as the execution time of the benchmarking program with stop-the-world gc time excluded.

As discussed in chapter 3, I implemented the Garbage-first family of collectors by performing progressive improvements on a simple region-based collector. In this way, after performing an algorithmic improvement on a collector, we can measure the overhead of the newly involved barriers or other technologies by comparing the benchmarking results of the old collector and the new collector.

As an output of the analysis of the overhead data, the average barrier overhead for each GC, each benchmark suite, each heap size and each barrier this project used is reported as well as their corresponding 95% confidence interval. The overall average overhead and 95% confidence interval for each barrier is also reported.

Overheads (ns)	Average	95% CI
637 M	177.74%	±314.43%
939 M	93.03%	±243.70%
1414 M	57.57%	±99.94%
1971 M	45.84%	±87.38%
Overall	92.93%	±231.72%

Table 4.4: Snapshot-at-the-begining barrier overhead

Overheads (ns)	Average	95% CI
637 M	58.69%	±206.84%
939 M	73.64%	±176.09%
1414 M	67.01%	±119.19%
1971 M	63.64%	±140.67%
Overall	65.81%	±164.54%

Table 4.5: Remembered set barrier overhead

4.5.2 Snapshot-at-the-begining barriers

Table 4.4 shows the overheads of the Snapshot-at-the-begining barrier as well as their 95% confidence interval on different heap sizes.

Discussion

4.5.3 Remembered set barriers

Table 4.5 shows the overheads of the Remembered set barrier as well as their 95% confidence interval on different heap sizes.

Discussion

4.5.4 Brooks indirection pointer barriers

Table 4.6 shows the overheads of the Brooks indirection pointer barrier as well as their 95% confidence interval on different heap sizes.

Discussion

Overheads (ns)	Average	95% CI
637 M	88.09%	±162.98%
939 M	98.35%	±116.13%
1414 M	82.94%	±86.65%
1971 M	77.86%	±86.32%
Overall	86.76%	±117.59%

Table 4.6: Brooks indirection pointer barrier overhead

4.6   Remembered-Set Size

4.6.1   Evaluation metrices

4.6.2   Results

4.6.3   Discussion

4.7   Summary

This chapter discusses the measurement methodology and results for evaluating the performance of the G1 family of garbage collectors. This includes the evaluation of GC pause times and overhead of several mutator barriers. Also the phenomenon revealed in the measurement results is carefully discussed.

On the one hand, based on the measurement results, we can see that linear scan based evacuation increses the work for each GC. Using concurrent marking, concurrent evacuation or remembered set based partial heap scanning can significantly reduces the pause time for each GC. On the other hand, using mutator barrier barriers can largely increase the mutator throughput reduction.

Also based on the measurement results, the generational G1 and Shenandoah GC shows some disappointing performance on GC pause time and mutator overheads respecively. However, the time scope of this project is limited. In this way, there are still much optimization job and other work to do in the future, which will be discussed in Chapter 5.

Heap Size = 637 MB						
Footprint	Committed Memory			Utilization		
	Min	Mean	Max	Min	Mean	Max
Garbage-First Non-generational	2.88%	8.30% ±6.09%	30.64%	0.08%	2.35% ±7.60%	27.97%
Garbage-First Generational	2.77%	8.69% ±6.68%	23.44%	0.09%	2.14% ±4.69%	19.70%

Heap Size = 939 MB						
Footprint	Committed Memory			Utilization		
	Min	Mean	Max	Min	Mean	Max
Garbage-First Non-generational	2.61%	8.98% ±9.08%	38.23%	0.07%	1.82% ±6.09%	21.53%
Garbage-First Generational	2.26%	9.77% ±7.63%	26.08%	0.08%	2.19% ±4.81%	16.98%

Heap Size = 1414 MB						
Footprint	Committed Memory			Utilization		
	Min	Mean	Max	Min	Mean	Max
Garbage-First Non-generational	2.36%	7.98% ±6.10%	36.76%	0.03%	1.11% ±3.67%	21.50%
Garbage-First Generational	3.18%	9.85% ±6.93%	23.62%	0.08%	2.00% ±4.02%	22.55%

Heap Size = 1971 MB						
Footprint	Committed Memory			Utilization		
	Min	Mean	Max	Min	Mean	Max
Garbage-First Non-generational	2.25%	7.98% ±6.02%	37.00%	0.03%	0.91% ±2.76%	20.85%
Garbage-First Generational	3.22%	10.15% ±6.76%	21.72%	0.08%	2.81% ±6.64%	22.00%

Overall						
Footprint	Committed Memory			Utilization		
	Min	Mean	Max	Min	Mean	Max
Garbage-First Non-generational	2.25%	8.36% ±6.99%	38.23%	0.03%	1.74% ±6.09%	27.97%
Garbage-First Generational	2.26%	9.56% ±7.09%	26.08%	0.08%	2.27% ±5.14%	22.55%

Table 4.7: Remembered set footprint

Conclusion

This thesis is aiming to identify and explore the underlying relationship of the G1 family of garbage collectors as well as analyze the performance impact caused by each structural components. The explorations and discussion in the previous chapters have successfully demostrates that the relationship among the G1 family of collectors exists and can impact the GC performance in both positive and negative ways. Hence the pros and cons of each algorithm and the cause of these phenomenons are discussed based on the evaluation results.

5.1   Future Work

Although this thesis has come to an end, the research on this topic is far to finish. There are still much work to do after this research project. In general, a few typical future works are: 1. Resolve a data rece problem for Shenandoah GC. 2. Perform some optimizations 3. Exploring C4 GC and ZGC. 4. Further G1 related GC research

5.1.1   Race problem for Shenandoah GC

Currently the implementation of Shenandoah GC has a data race problem. Based on the previous analysis, the mutator is not lock and release a monitor expectedly during the concurrent phase of the Shenandoah GC, when there can be two different copies of an object exist in the heap.

Duo to the time scope of thie thesis, I cannot solve this issue currently. But currently a work-around is implemented in the Shenandoah GC, with a lower mutator performance. This explains the bad performance of mutator overhead we evaluated in chapter 4. As the most urgent problem I are facing, I plan to solve this issue immediately after this project.

5.1.2   Optimizations

As part of the mutator latency results discussed in chapter 4, the generational G1 currently does not reveal any peprformance gain compared to the non-generational G1. It is expected to have more optimizations and parameter tuning to the generational G1 to further increase its performance.

I am also considering the possibility to make the implemented Garbage-first and Shenandoah GC become production ready. One major part missing is the Optimizations since the development of these garbage collectors was following the original design of these collectors and did not have too much optimizations. After performing some optimizations on these collectors as well as some additional correctness verfication, the collector can have the possibility to become production ready.

5.1.3   C4 GC and ZGC

Due to the time scope of this project, I did not implement and measure C4 GC and ZGC. In addition, since JikesRVM and MMTk only supports the 32bit address space but the pointer coloring process in ZGC requires the 64bit address space, this makes the implementation of ZGC more difficult.

However, as the most latest member of the Garbage-first family of collectors, C4/ZGC generally has better performance than G1 GC and Shenandoah GC. By performing most of the GC work in concurrent, the pause time of C4/ZGC is not propotional to the heap size and are expected to be less than 10 milliseconds even targeting 100 GB heaps (Liden and Karlsson [2018]). In this way these two collectors are extremely worth for an exploration.

However, a detailed plan of resolving several hardware incompactabilities should be done in the future, before starting the implementation of these two collectors.

5.1.4   Future G1 related GC research

The Garbagr-first family of garbage collectors have been proved to be high performance, in terms of GC latency and mutator overheads. However, the measurement results is still not perfect, which means there are a lot more can be done to further improve the GC performance.

One example can be adding a generational extension to the Shenandoah GC to collect young garbage as early as possible to reduce the frequency of falling to full GCs. This generational mode involves a remembered set to remember object pointers in mature space pointing to the nursery space. However, in addition to using the table-based remembered sets used in G1, the buffer-based remembered sets introduced by Blackburn and McKinley [2008] can also be explored as a comparison to G1's remembered sets.

There are still many details of G1 related garbage collection algorithms can be explored. By performing more improvements on the existing Garbagr-first family of collectors (e.g. the Shenandoah GC), the GC performance can be expected to have more improvements.

5.2   Summary

This thesis is aiming to identify and explore the underlying relationship of the G1 family of garbage collectors, implement them as a series of collectors to reflect such

relationships and analyze the GC performance impact of different algorithm components.

As discussed in previous chapters, the potential relationships among the G1 family of collectors are successfully discovered and discussed. Based on such relationships, this thesis produces the first implementation of the G1 family of collectors that reflect to the underlying algorithmic relationships. Most of my implementations result in a reasonable or even unexpectedly performance for either GC pause time and mutator overheads. Based on these implementations and measurement results, the pros and cons as well as their underlying reasons are explored and discussed. Hence, the explorations performed in this thesis can inspire GC designers to reconsider the design and structure of region-based GC algorithms to make further valuable improvements and invent new algorithms.

In conclusion, the relationships among the G1 family of collectors exist and have impact to the GC performance in many ways, including negative performance impacts. Which mean there is still a lot to improve and more research should be done in this area.

Bibliography

ALPERN, B.; AUGART, S.; BLACKBURN, S. M.; BUTRICO, M.; COCCHI, A.; CHENG, P.;
DOLBY, J.; FINK, S.; GROVE, D.; HIND, M.; ET AL., 2005. The jikes research virtual
machine project: building an open-source research community. *IBM Systems Journal*,
44, 2 (2005), 399–417. (cited on page 5)

BLACKBURN, S. M.; CHENG, P.; AND MCKINLEY, K. S., 2004. Oil and water? high
performance garbage collection in java with mmtk. In *Software Engineering, 2004.
ICSE 2004. Proceedings. 26th International Conference on*, 137–146. IEEE. (cited on
page 3)

BLACKBURN, S. M.; GARNER, R.; HOFFMANN, C.; KHANG, A. M.; MCKINLEY, K. S.;
BENTZUR, R.; DIWAN, A.; FEINBERG, D.; FRAMPTON, D.; GUYER, S. Z.; HIRZEL, M.;
HOSKING, A.; JUMP, M.; LEE, H.; MOSS, J. E. B.; PHANSALKAR, A.; STEFANOVIĆ, D.;
VANDRUNEN, T.; VON DINCKLAGE, D.; AND WIEDERMANN, B., 2006. The dacapo
benchmarks: Java benchmarking development and analysis. *SIGPLAN Not.*, 41,
10 (Oct. 2006), 169–190. doi:10.1145/1167515.1167488. http://doi.acm.org/10.1145/
1167515.1167488. (cited on page 35)

BLACKBURN, S. M. AND HOSKING, A. L., 2004. Barriers: Friend or foe? In *Proceedings
of the 4th international symposium on Memory management*, 143–151. ACM. (cited on
page 37)

BLACKBURN, S. M. AND MCKINLEY, K. S., 2008. Immix: a mark-region garbage collector
with space efficiency, fast collection, and mutator performance. In *ACM SIGPLAN
Notices*, vol. 43, 22–32. ACM. (cited on page 46)

BRIANDEMSKY, K. L. G., 2015. Speculative region-based memory management for big
data systems. (2015). (cited on page 2)

BROOKS, R. A., 1984. Trading data space for reduced time and code space in real-time
garbage collection on stock hardware. In *Proceedings of the 1984 ACM Symposium on
LISP and functional programming*, 256–262. ACM. (cited on page 14)

DETLEFS, D.; FLOOD, C.; HELLER, S.; AND PRINTEZIS, T., 2004. Garbage-first garbage
collection. In *Proceedings of the 4th international symposium on Memory management*,
37–48. ACM. (cited on pages 1, 8, 10, 14, 21, 24, and 26)

DETLEFS, D. L.; MARTIN, P. A.; MOIR, M.; AND STEELE JR, G. L., 2002. Lock-free
reference counting. *Distributed Computing*, 15, 4 (2002), 255–271. (cited on page 6)

Endo, T.; Taura, K.; and Yonezawa, A., 1997. A scalable mark-sweep garbage collector on large-scale shared-memory machines. In *Proceedings of the 1997 ACM/IEEE conference on Supercomputing*, 1–14. ACM. (cited on page 7)

Flood, C. H.; Kennke, R.; Dinn, A.; Haley, A.; and Westrelin, R., 2016. Shenandoah: An open-source concurrent compacting garbage collector for openjdk. In *Proceedings of the 13th International Conference on Principles and Practices of Programming on the Java Platform: Virtual Machines, Languages, and Tools*, 13. ACM. (cited on pages 1, 10, 14, and 28)

Gay, D. and Aiken, A., 1998. *Memory management with explicit regions*, vol. 33. ACM. (cited on page 10)

Liden, P. and Karlsson, S., 2018. Zgc: A scalable low-latency garbage collector. http://openjdk.java.net/jeps/333. (cited on pages 1, 9, 14, and 46)

Lins, R. D., 1992. Cyclic reference counting with lazy mark-scan. *Information Processing Letters*, 44, 4 (1992), 215–220. (cited on page 7)

Qian, F. and Hendren, L., 2002. An adaptive, region-based allocator for java. In *ACM SIGPLAN Notices*, vol. 38, 127–138. ACM. (cited on page 10)

Tene, G.; Iyengar, B.; and Wolf, M., 2011. C4: The continuously concurrent compacting collector. *ACM SIGPLAN Notices*, 46, 11 (2011), 79–88. (cited on pages 1, 9, and 14)

Yang, X.; Blackburn, S. M.; Frampton, D.; and Hosking, A. L., 2012. Barriers reconsidered, friendlier still! In *ACM SIGPLAN Notices*, vol. 47, 37–48. ACM. (cited on pages 11 and 37)

Yuasa, T., 1990. Real-time garbage collection on general-purpose machines. *Journal of Systems and Software*, 11, 3 (1990), 181–198. (cited on pages 14, 15, and 20)